

Discovering Groups

Summary

Being able to meaningfully cluster data into groups using Clustering or Cluster Analysis is a key part of the process of exploratory and descriptive data mining. Clustering techniques are all a form of unsupervised machine learning. Numerous techniques for clustering exist; in this lecture we'll look at two of the most common and useful: Hierarchical Clustering and K-Means Clustering. We'll also look briefly at a more advanced, but computationally intensive algorithm called Mean Shift Clustering that both produces clusters and finds the *modes* of the data.

Key points

Clustering

- Clustering is an unsupervised machine learning technique, that learns to group data without prior knowledge of what the groups should look like.

Hierarchical Clustering

- Hierarchical Clustering attempts to iteratively break data into a hierarchy of clusters
- Hierarchical Agglomerative Clustering builds a binary tree of clusters from the leaf nodes upwards towards the root
 - Known as a bottom-up approach
 - Requires three things:
 - a set of items to cluster
 - a distance measure to measure how close items are to each other
 - e.g. an L_p distance, a similarity measure converted to a distance (i.e. 1-Pearson or 1-cosine)
 - Doesn't necessarily have to be a distance computed over a vector; some forms of agglomerative clustering allow only need a matrix of distances or similarities computed between all items as input (see below)
 - a linkage criterion which measures dissimilarity of clusters as a function of the pairwise distances of items in the clusters
 - Basic approach:

>-

 - Initially every item is in a cluster of its own
 - While there is more than one single cluster:
 - The closest pair of clusters according to the linkage criterion are merged into a bigger cluster
 - By recording the merges at each step a binary tree structure linking the clusters can be formed
 - Often a useful way of utilising this is by drawing a diagram known as a dendrogram that shows the structure of the tree
- Two categories of linkage criterion:

- Centroid-based linkage functions that measure similarity between clusters based on the distance between their centroids
 - Requires that each item is represented by a numeric feature vector that can be interpreted as a position in space
 - Examples:
 - Weighted Centroid Clustering (*WPGMC – Weighted Pair Group Method with Centroids; often also known as the “median” method*)
 - When two clusters s and t are combined into a new cluster u , the average of centroids s and t give the new centroid u
 - Unweighted Centroid Clustering (*UPGMC – Unweighted Pair Group Method with Centroids*)
 - When two clusters s and t are combined into a new cluster u , the average of the positions of all the items within s and t give the new centroid u
- Distance-based linkage functions that measure distances between clusters as a function of the distances between items within those clusters.
 - Clustering can be performed purely as a function of a distance matrix in which each element $D_{i,j}$ represents the distance, $d(i,j)$, between items i and j
 - Commonly used linkage criteria between two sets (clusters) of items A and B include:
 - Minimum or single-linkage clustering: $\min\{d(a,b) : a \in A, b \in B\}$
 - Drawback: tends to produce long, thin, clusters where the items at each end are far apart
 - Maximum or complete-linkage clustering: $\max\{d(a,b) : a \in A, b \in B\}$
 - Avoids problems of single-linkage clustering; tends to find compact clusters of approximately equal diameter
 - Mean or average linkage clustering (UPGMA – Unweighted Pairwise Group Method with Arithmetic Mean):

$$\frac{1}{|A||B|} \sum_{a \in A} \sum_{b \in B} d(a,b)$$
- In general, complexity is $O(n^3)$, which can be a problem for large data sets, however there are some $O(n^2)$ variants for the single-linkage and complete-linkage cases
- Divisive clustering algorithms (“top-down” approaches), which start with all the data in the root node and recursively split do exist
 - Not widely used in practice.
 - One major reason is that in general complexity is $O(2^n)$, which is worse than the agglomerative methods.

K-Means Clustering

- The K-Means algorithm (also known as *Lloyds algorithm*) is a simple, but powerful, approach to clustering that attempts to group data in a feature space into K groups or clusters represented by centroids (i.e. the mean point of the class in feature-space).
- Algorithm:



- The K -value must be chosen *a-priori* (beforehand)
- To begin, K initial cluster centres are chosen (typically randomly or from a sample of the existing data points, although note that better initialisation procedures exist - e.g. the KMeans++ algorithm)
- Then the following process is performed iteratively until the centroids don't move between iterations (or the maximum number of iterations is reached):

- Each point is assigned to its closest centroid
 - The centroid is recomputed as the mean of all the points assigned to it. If the centroid has no points assigned it is randomly re-initialised to a new point.
- The final clusters are created by assigning all points to their nearest centroid.
- K-Means always converges, but not necessarily to the most optimal solution

Mean Shift Clustering

- Mean Shift is a standard algorithm to efficiently find the modes of a Probability Density Function (PDF) from a set of samples of that PDF (i.e. the feature vectors representing a set of items).
 - The only variable of the mean shift algorithm is the kernel and the kernel bandwidth of a kernel density estimator.
 - Clustering is an application of the mean shift procedure
 - Automatically chooses the number of clusters!
- The PDF of a continuous random variable is a function that describes the relative likelihood for this random variable to take on a given value
 - The PDF is non-negative everywhere and sums to 1
- In the context of a feature space, the PDF is a function that tells you how likely it is that a feature vector is *drawn* from a specific location in a feature space.
 - A feature vector drawn from part of the space where there are lots of similar items would have a higher probability density than if the drawn feature vector were from a part of the space with very few similar items
 - or in other words, dense parts of the space with more items have a higher probability density
 - Generally speaking, for arbitrary features describing a set of items, the PDF cannot be described empirically
 - Must be estimated using some other method
 - Simple, but crude, way to do this would be to quantise the feature space into bins in order to build a histogram
 - Each bin would contain the count of the number of items with feature vectors falling into that bin divided by the number of total items
 - Major disadvantage of this approach is that it isn't *continuous* and only gives a discrete approximation of the PDF
 - Better way to do this is to use a Kernel Density Estimator (also known as a "Parzen Window")
 - Letting (x_1, x_2, \dots, x_n) represent the set of samples (e.g. feature vectors) in a d -dimensional space R^d from an unknown density f , then:

$$f(\mathbf{x}) = \frac{1}{nh^d} \sum_{i=1}^n K\left(\frac{\mathbf{x} - \mathbf{x}_i}{h}\right)$$

where $K(\cdot)$ is the kernel (a non-negative function that integrates to one and has mean zero), and $h > 0$ is a smoothing parameter called the kernel bandwidth.

- Common choice for the kernel is a multivariate Gaussian with zero mean and unit s.d.
 - For radially symmetric kernels, it suffices to define the profile of the kernel $k(x)$ satisfying $K(x) = c_{k,d} k(\|x\|^2)$
- Intuitively one wants to choose h as small as the data will allow
 - there is always a trade-off between the bias of the estimator and its variance however

- The Mean Shift procedure attempts to find the modes of the density function – that is the points where the gradient is 0: $\nabla f(\mathbf{x})=0$

- Assuming a radially symmetric kernel, then the gradient is:

$$\begin{aligned}\nabla f(\mathbf{x}) &= \frac{2c_{k,d}}{nh^{d+2}} \sum_{i=1}^n (\mathbf{x} - \mathbf{x}_i) g\left(\left\|\frac{\mathbf{x} - \mathbf{x}_i}{h}\right\|^2\right) \\ &= \frac{2c_{k,d}}{nh^{d+2}} \left[\sum_{i=1}^n g\left(\left\|\frac{\mathbf{x} - \mathbf{x}_i}{h}\right\|^2\right) \right] \left[\frac{\sum_{i=1}^n \mathbf{x}_i g\left(\left\|\frac{\mathbf{x} - \mathbf{x}_i}{h}\right\|^2\right)}{\sum_{i=1}^n g\left(\left\|\frac{\mathbf{x} - \mathbf{x}_i}{h}\right\|^2\right)} - \mathbf{x} \right]\end{aligned}$$

where $g(s) = -k'(s)$.

- The first term in the above is proportional to the density estimate at \mathbf{x} computed with a kernel $G(\mathbf{x}) = c_{g,d} g(\|\mathbf{x}\|^2)$, and the second term

$$\mathbf{m}_h(\mathbf{x}) = \frac{\sum_{i=1}^n \mathbf{x}_i g\left(\left\|\frac{\mathbf{x} - \mathbf{x}_i}{h}\right\|^2\right)}{\sum_{i=1}^n g\left(\left\|\frac{\mathbf{x} - \mathbf{x}_i}{h}\right\|^2\right)} - \mathbf{x}$$

is the mean shift.

- The mean shift vector always points toward the direction of the maximum increase in the density.
- The mean shift procedure, obtained by successive
 - computation of the mean shift vector $\mathbf{m}_h(\mathbf{x}_t)$,
 - translation of the window $\mathbf{x}_{t+1} = \mathbf{x}_t + \mathbf{m}_h(\mathbf{x}_t)$

is guaranteed to converge to a point where the gradient of density function is zero

- Mean Shift Clustering works as follows:



- for each feature vector:
 - apply the mean shift procedure until convergence and store the resultant mode
- the set of feature vectors that converge to the same mode define the *basin of attraction* of that mode; all features that converged to the same mode belong to the same cluster

Further Reading

- Chapter 3 of “Programming Collective Intelligence” gives a good overview of the some of the basic techniques.
- Relevant sections of Chapter 14 of The Elements of Statistical Learning (https://web.stanford.edu/~hastie/ElemStatLearn/printings/ESLII_print10.pdf) provide a good academic introduction
- Wikipedia has reasonable commentary (and good links to the original research) on a number of the topics:
 - https://en.wikipedia.org/wiki/Hierarchical_clustering (https://en.wikipedia.org/wiki/Hierarchical_clustering)
 - https://en.wikipedia.org/wiki/K-means_clustering (https://en.wikipedia.org/wiki/K-means_clustering)
 - https://en.wikipedia.org/wiki/Mean_shift (https://en.wikipedia.org/wiki/Mean_shift)
- k-means++: the advantages of careful seeding (<http://ilpubs.stanford.edu:8090/778/1/2006-13.pdf>) . Arthur and Vassilvitskii. Proceedings of the eighteenth annual ACM–SIAM symposium on Discrete algorithms. Society for Industrial and Applied Mathematics Philadelphia, PA, USA. pp. 1027–1035. 2007.
- Mean shift: A robust approach toward feature space analysis (<http://www.caip.rutgers.edu/riul/research/papers/pdf/mnshft.pdf>) . Comaniciu and Meer. IEEE Trans. Pattern Anal. Machine Intell., 24:603–619, 2002.
- Good descriptions of clustering techniques:

- Learning from Data: Concepts, Theory, and Methods (2nd ed.). Cherkassky and Mulier. John Wiley & Sons, Inc., New York, NY, USA.